

# **BCG Task 2**

**Kiana kheiri**

## **Exploratory Data Analysis&Data Cleaning**

## The Datasets

The dataset `ml_case_training_output.csv` named as `pco_output` contains:

- `id`: contact id \* `churn`: has the client churned over the next 3 months

The dataset `ml_case_training_hist_data.csv` named as `pco_hist` contains the history of energy and power consumption per client:

- `id`: contact id *`price_date`: reference date* `price_p1_var`: price of energy for the 1st period *`price_p2_var`: price of energy for the 2nd period* `price_p3_var`: price of energy for the 3rd period *`price_p1_fix`: price of power for the 1st period* `price_p2_fix`: price of power for the 2nd period \* `price_p3_fix`: price of power for the 3rd period

The dataset `ml_case_training_data.csv` contains:

- `id` : contact id
- `activity_new` : category of the company's activity. 419 unique values, remove NaN
- `campaign_disc_elec` : code of the electricity campaign the customer last subscribed to. 0 non-null
- `channel_sales` : code of the sales channel
- `cons_12m` : electricity consumption of the past 12 months
- `cons_gas_12m` : gas consumption of the past 12 months
- `cons_last_month` : electricity consumption of the last month
- `date_activ` : date of activation of the contract
- `date_end` : registered date of the end of the contract
- `date_first_activ` : date of first contract of the client
- `date_modif_prod` : date of last modification of the product
- `date_renewal` : date of the next contract renewal
- `forecast_base_bill_ele` : forecasted electricity bill baseline for next month
- `forecast_base_bill_year` : forecasted electricity bill baseline for calendar year
- `forecast_bill_12m` : forecasted electricity bill baseline for 12 months
- `forecast_cons` : forecasted electricity consumption for next month
- `forecast_cons_12m` : forecasted electricity consumption for next 12 months
- `forecast_cons_year` : forecasted electricity consumption for next calendar year
- `forecast_discount_energy` : forecasted value of current discount
- `forecast_meter_rent_12m` : forecasted bill of meter rental for the next 12 months
- `forecast_price_energy_p1` : forecasted energy price for 1st period
- `forecast_price_energy_p2` : forecasted energy price for 2nd period

## LIBRARIES

### Gathering data

```
In [1]: #Import packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set(color_codes=True)
import pickle
```

```
In [2]: #Loading data
train_data=pd.read_csv('ml_case_training_data.csv')
history_data=pd.read_csv('ml_case_training_hist_data.csv')
churn_data=pd.read_csv('ml_case_training_output.csv')
```

```
In [3]: #Show the first 5 rows of data
train_data.head()
```

Out[3]:

	id	activity_new	campaign_disc_ele
0	48ada52261e7cf58715202705a0451c9	esoiiifxdbkcsluxmfuacbdckommixw	NaN Imket
1	24011ae4ebbe3035111d65fa7c15bc57	NaN	NaN foo
2	d29c2c54acc38ff3c0614d0a653813dd	NaN	NaN
3	764c75f661154dac3a6c254cd082ea7d	NaN	NaN foo
4	bba03439a292a1e166f80264c16191cb	NaN	NaN Imket

5 rows × 32 columns

```
In [4]: history_data.head()
```

Out[4]:

	id	price_date	price_p1_var	price_p2_var	price_p3_var	price
0	038af19179925da21a25619c5a24b745	2015-01-01	0.151367	0.0	0.0	44.1
1	038af19179925da21a25619c5a24b745	2015-02-01	0.151367	0.0	0.0	44.1
2	038af19179925da21a25619c5a24b745	2015-03-01	0.151367	0.0	0.0	44.1
3	038af19179925da21a25619c5a24b745	2015-04-01	0.149626	0.0	0.0	44.1
4	038af19179925da21a25619c5a24b745	2015-05-01	0.149626	0.0	0.0	44.1

```
In [5]: churn_data.head()
```

Out[5]:

	id	churn
0	48ada52261e7cf58715202705a0451c9	0
1	24011ae4ebbe3035111d65fa7c15bc57	1
2	d29c2c54acc38ff3c0614d0a653813dd	0
3	764c75f661154dac3a6c254cd082ea7d	0
4	bba03439a292a1e166f80264c16191cb	0

```
In [6]: #merge the train_data and churn_data into one dataframe
train=pd.merge(train_data,churn_data, on="id")
train.head()
```

Out[6]:

	id	activity_new	campaign_disc_ele
0	48ada52261e7cf58715202705a0451c9	esoiifxdlbkcsluxmfuacbdckommixw	NaN Imket
1	24011ae4ebbe3035111d65fa7c15bc57	NaN	NaN foo:
2	d29c2c54acc38ff3c0614d0a653813dd	NaN	NaN
3	764c75f661154dac3a6c254cd082ea7d	NaN	NaN foo:
4	bba03439a292a1e166f80264c16191cb	NaN	NaN Imket

5 rows × 33 columns

Accessing Data

```
In [7]: #See the datatype of train data
train.dtypes
```

```
Out[7]: id                object
activity_new             object
campaign_disc_ele        float64
channel_sales            object
cons_12m                 int64
cons_gas_12m             int64
cons_last_month          int64
date_activ               object
date_end                 object
date_first_activ         object
date_modif_prod          object
date_renewal             object
forecast_base_bill_ele   float64
forecast_base_bill_year  float64
forecast_bill_12m        float64
forecast_cons            float64
forecast_cons_12m        float64
forecast_cons_year       int64
forecast_discount_energy float64
forecast_meter_rent_12m  float64
forecast_price_energy_p1 float64
forecast_price_energy_p2 float64
forecast_price_pow_p1    float64
has_gas                  object
imp_cons                 float64
margin_gross_pow_ele     float64
margin_net_pow_ele       float64
nb_prod_act              int64
net_margin               float64
num_years_antig          int64
origin_up                object
pow_max                  float64
churn                    int64
dtype: object
```

```
In [8]: history_data.dtypes
```

```
Out[8]: id                object
price_date               object
price_p1_var             float64
price_p2_var             float64
price_p3_var             float64
price_p1_fix             float64
price_p2_fix             float64
price_p3_fix             float64
dtype: object
```

```
In [9]: #See the shape of dataset
train.shape
```

```
Out[9]: (16096, 33)
```

```
In [10]: history_data.shape
```

```
Out[10]: (193002, 8)
```

```
In [11]: #See the general descriptive statistics of data
train.describe()
```

```
Out[11]:
```

	campaign_disc_ele	cons_12m	cons_gas_12m	cons_last_month	forecast_base_bill_ele
<b>count</b>	0.0	1.609600e+04	1.609600e+04	1.609600e+04	3508.000000
<b>mean</b>	NaN	1.948044e+05	3.191164e+04	1.946154e+04	335.843857
<b>std</b>	NaN	6.795151e+05	1.775885e+05	8.235676e+04	649.406000
<b>min</b>	NaN	-1.252760e+05	-3.037000e+03	-9.138600e+04	-364.940000
<b>25%</b>	NaN	5.906250e+03	0.000000e+00	0.000000e+00	0.000000
<b>50%</b>	NaN	1.533250e+04	0.000000e+00	9.010000e+02	162.955000
<b>75%</b>	NaN	5.022150e+04	0.000000e+00	4.127000e+03	396.185000
<b>max</b>	NaN	1.609711e+07	4.188440e+06	4.538720e+06	12566.080000

8 rows × 23 columns

It's seems that the campaign\_disc\_lal is an empty column

```
In [12]: history_data.describe()
```

```
Out[12]:
```

	price_p1_var	price_p2_var	price_p3_var	price_p1_fix	price_p2_fix	price_p3_fix
<b>count</b>	191643.000000	191643.000000	191643.000000	191643.000000	191643.000000	191643.000000
<b>mean</b>	0.140991	0.054412	0.030712	43.325546	10.698201	6.45543
<b>std</b>	0.025117	0.050033	0.036335	5.437952	12.856046	7.78227
<b>min</b>	0.000000	0.000000	0.000000	-0.177779	-0.097752	-0.06517
<b>25%</b>	0.125976	0.000000	0.000000	40.728885	0.000000	0.00000
<b>50%</b>	0.146033	0.085483	0.000000	44.266930	0.000000	0.00000
<b>75%</b>	0.151635	0.101780	0.072558	44.444710	24.339581	16.22638
<b>max</b>	0.280700	0.229788	0.114102	59.444710	36.490692	17.45822

```
In [13]: #See The missing data of train
train.isnull().sum()/train.shape[0]
```

```
Out[13]: id                                0.000000
activity_new                             0.593004
campaign_disc_ele                        1.000000
channel_sales                           0.262053
cons_12m                                0.000000
cons_gas_12m                            0.000000
cons_last_month                         0.000000
date_activ                              0.000000
date_end                                0.000124
date_first_activ                        0.782058
date_modif_prod                         0.009754
date_renewal                            0.002485
forecast_base_bill_ele                  0.782058
forecast_base_bill_year                  0.782058
forecast_bill_12m                       0.782058
forecast_cons                           0.782058
forecast_cons_12m                       0.000000
forecast_cons_year                      0.000000
forecast_discount_energy                 0.007828
forecast_meter_rent_12m                  0.000000
forecast_price_energy_p1                 0.007828
forecast_price_energy_p2                 0.007828
forecast_price_pow_p1                    0.007828
has_gas                                 0.000000
imp_cons                                0.000000
margin_gross_pow_ele                     0.000808
margin_net_pow_ele                       0.000808
nb_prod_act                             0.000000
net_margin                              0.000932
num_years_antig                          0.000000
origin_up                               0.005405
pow_max                                  0.000186
churn                                    0.000000
dtype: float64
```

As we can see that some of columns have missing data over 50%, we need to clean them in the later

```
In [14]: history_data.isnull().sum()/history_data.shape[0]
```


```
Out[14]: id                                0.000000
price_date                               0.000000
price_p1_var                             0.007041
price_p2_var                             0.007041
price_p3_var                             0.007041
price_p1_fix                             0.007041
price_p2_fix                             0.007041
price_p3_fix                             0.007041
dtype: float64
```



```
In [15]: #Deep diving on the main parameters,first for the Churn
churn=train[['id','churn']]
churn.columns=['Companies','churn']
```

```
In [16]: churn_total=churn.groupby(churn['churn']).count()
churn_percentage=churn_total/churn_total.sum()*100
```


```
In [17]: ax=churn_percentage.transpose().plot(kind='bar',stacked=True,figsize=(8,
6),rot=0)
for p in ax.patches:
    value=str(round(p.get_height(),1))
    if value=='0':
        continue
    ax.annotate(value,((p.get_x()+p.get_width()/2)*0.5,p.get_y()+p.get_h
eight()/2*0.6),
                color='white',size=(8))
plt.legend(['Retention','Churn'],loc="upper right")
plt.ylabel("Company base(%)");
```



About 10% of total customers have churned

```
In [18]: #Next see the acitivity distribution
activity=train[['id','activity_new','churn']]
activity=activity.groupby([activity['activity_new'],activity['churn']])[
'id'].count().unstack(level=1).sort_values(by=[0],ascending=False)
```

```
In [19]: activity.plot(kind='bar',figsize=(12,10),width=2,stacked=True,title="SME  
Activity")  
plt.ylabel("Number of Companies")  
plt.xlabel('Activity')  
plt.legend(['Retention','Churn'],loc="upper right")  
plt.xticks([])  
plt.show()
```



The xticks is not showing to facilitate the visualization and the distribution of the activity is despite the lack of 60% of the entries

```
In [20]: activity_total=activity.fillna(0)[0]+activity.fillna(0)[1]
activity_percentage=activity.fillna(0)[1]/(activity_total)*100
pd.DataFrame({'Percentage churn':activity_percentage,
              'Total companies':activity_total}).sort_values(by='Percentage churn',ascending=False).head()
```


Out[20]:

	Percentage churn	Total companies
activity_new		
xwkaesbkfsacseixksofpddwfkbobki	100.0	1.0
wkwddcuiboaeaalcaawlwmldiwpewma	100.0	1.0
ikiucmkuisupefxcfxkulpwssppfu	100.0	1.0
opoiuudmxdssidluooopfswlkkcsxf	100.0	1.0
pfcocskbxlmofswiflsbcefcupfbopuo	100.0	2.0

```
In [21]: #Now is about Sales channel
channel=train[['id','channel_sales','churn']]
channel=channel.groupby([channel['channel_sales'],channel['churn']])['id'].count().unstack(level=1).fillna(0)
```

```
In [22]: channel_churn=(channel.div(channel.sum(axis=1),axis=0)*100).sort_values(
by=[1],ascending=False)
```

```
In [23]: ax=channel_churn.plot(kind='bar',stacked=True,figsize=(8,6),rot=45)
for p in ax.patches:
    value=str(round(p.get_height(),1))
    if value=='0':
        continue
    ax.annotate(value,((p.get_x()+p.get_width()/2)*0.94,p.get_y()+p.get_
height()/2*0.99),
                color='white',size=(9))
plt.title('Sales Channel')
plt.legend(['Retention','Churn'],loc="upper right")
plt.ylabel("Company base(%)");
```



```
In [24]: channel_total=channel.fillna(0)[0]+channel.fillna(0)[1]
channel_percentage=channel.fillna(0)[1]/(channel_total)*100
pd.DataFrame({"Churn percentage":channel_percentage,
              "Total companies":channel_total}).sort_values(by='Churn per
centage',ascending=False).head()
```

Out[24]:


	Churn percentage	Total companies
channel_sales		
foosdfpfkusacimwksosbicdxkicaua	12.498306	7377.0
usilxuppasemubllopkaafesmlibmsdf	10.387812	1444.0
ewpakwlliwisiwduibdlfmalxowmwpci	8.488613	966.0
lmkebamcaaclubfxadlmueccxoimlema	5.595755	2073.0
epumfxlbckeskwexbiuasklxalciuu	0.000000	4.0

```
In [25]: #Next is the consumption
consumption=train[['id','cons_12m','cons_gas_12m','cons_last_month','imp
_cons','has_gas','churn']]
```

```

In [26]: fig,axs=plt.subplots(nrows=4,figsize=(8,15))
cons_12m=pd.DataFrame({'Retention':consumption[consumption['churn']==0][
'cons_12m'],
                        'Churn':consumption[consumption['churn']==1]['cons
_12m']})
cons_12m[['Retention','Churn']].plot(kind='hist',bins=20,ax=axs[0],stack
ed=True);
axs[0].set_xlabel('cons_12m')
axs[0].ticklabel_format(style='plain',axis='x')
cons_gas_12m=pd.DataFrame({'Retention':consumption[consumption['has_gas'
]=='t'][consumption[consumption['has_gas']=='t']['churn']==0]['cons_gas_
12m'],
                        'Churn':consumption[consumption['has_gas']=='t'][c
onsumption[consumption['has_gas']=='t']['churn']==1]['cons_gas_12m']})
cons_gas_12m[['Retention','Churn']].plot(kind='hist',bins=20,ax=axs[1],s
tacked=True);
axs[1].set_xlabel('cons_12m')
axs[1].ticklabel_format(style='plain',axis='x')
cons_last_month=pd.DataFrame({'Retention':consumption[consumption['chur
n']==0]['cons_last_month'],
                        'Churn':consumption[consumption['churn']==1]['cons
_last_month']})
cons_last_month[['Retention','Churn']].plot(kind='hist',bins=20,ax=axs[2
],stacked=True);
axs[2].set_xlabel('cons_12m')
axs[2].ticklabel_format(style='plain',axis='x')
imp_cons=pd.DataFrame({'Retention':consumption[consumption['churn']==0][
'imp_cons'],
                        'Churn':consumption[consumption['churn']==1]['imp_
cons']})
imp_cons[['Retention','Churn']].plot(kind='hist',bins=20,ax=axs[3],stack
ed=True);
axs[3].set_xlabel('imp_cons')
axs[3].ticklabel_format(style='plain',axis='x')

```




The distribution of the consumptions is highly right skewed and has a long tail, we need to check the outliers by use boxplot

```
In [27]: fig,axs=plt.subplots(nrows=4,figsize=(10,15))
sns.boxplot(consumption['cons_12m'],ax=axs[0])
sns.boxplot(consumption[consumption['has_gas']=='t']['cons_gas_12m'],ax=
axs[1])
sns.boxplot(consumption['cons_last_month'],ax=axs[2])
sns.boxplot(consumption['imp_cons'],ax=axs[3])
for ax in axs:
    ax.ticklabel_format(style='plain',axis='x')
axs[0].set_xlim(-200000,2000000)
axs[1].set_xlim(-200000,2000000)
axs[2].set_xlim(-20000,100000)
plt.show();
```



```
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```



It clearly that we can see the outliers and we will deal with them in the data cleaning

```
In [28]: #Now is about Dates
dates=train[['id','date_activ','date_end','date_modif_prod','date_renewal',
            'churn']].copy()
dates['date_activ']=pd.to_datetime(dates['date_activ'],format='%Y-%m-%d')
dates['date_end']=pd.to_datetime(dates['date_end'],format='%Y-%m-%d')
dates['date_modif_prod']=pd.to_datetime(dates['date_modif_prod'],format=
'%Y-%m-%d')
dates['date_renewal']=pd.to_datetime(dates['date_renewal'],format='%Y-%m-%d')
```

```
In [29]: def line_format(label):
        """
        Convert time label to the format of pandas line plot
        """
        month=label.month_name()[ :1]
        if label.month_name()=="January":
            month+=f'\n{label.year}'
        return month
```

```

In [30]: fig,axs=plt.subplots(nrows=4,figsize=(18,15))
date_activ=dates[['date_activ','churn','id']].set_index('date_activ').groupby([pd.Grouper(freq='M'),'churn']).count().unstack(level=1)
date_activ.plot(kind='bar',stacked=True,rot=0,ax=axs[0])
axs[0].set_xticklabels(map(lambda x:line_format(x),date_activ.index),fontsize=8)
axs[0].set_ylabel("Number of companies")
axs[0].legend(['Retention','Churn'],loc='upper right')
date_end=dates[['date_end','churn','id']].set_index('date_end').groupby([pd.Grouper(freq='M'),'churn']).count().unstack(level=1)
date_end.plot(kind='bar',stacked=True,rot=0,ax=axs[1])
axs[1].set_xticklabels(map(lambda x:line_format(x),date_end.index),fontsize=8)
axs[1].set_ylabel("Number of companies")
axs[1].legend(['Retention','Churn'],loc='upper right')
date_modif_prod=dates[['date_modif_prod','churn','id']].set_index('date_modif_prod').groupby([pd.Grouper(freq='M'),'churn']).count().unstack(level=1)
date_modif_prod.plot(kind='bar',stacked=True,rot=0,ax=axs[2])
axs[2].set_xticklabels(map(lambda x:line_format(x),date_modif_prod.index),fontsize=8)
axs[2].set_ylabel("Number of companies")
axs[2].legend(['Retention','Churn'],loc='upper right')
date_renewal=dates[['date_renewal','churn','id']].set_index('date_renewal').groupby([pd.Grouper(freq='M'),'churn']).count().unstack(level=1)
date_renewal.plot(kind='bar',stacked=True,rot=0,ax=axs[3])
axs[3].set_xticklabels(map(lambda x:line_format(x),date_renewal.index),fontsize=8)
axs[3].set_ylabel("Number of companies")
axs[3].legend(['Retention','Churn'],loc='upper right');

```



However, the date's distribution seems does not provide any insight.

```
In [31]: #Now is about the forecast
forecast=train[['id', 'forecast_base_bill_ele', 'forecast_base_bill_year',
'forecast_bill_12m', 'forecast_cons',
'forecast_cons_12m', 'forecast_cons_year', 'forecast_discount_energy',
'forecast_meter_rent_12m', 'forecast_price_energy_p1', 'forecast_price_energy_p2',
'forecast_price_pow_p1', 'churn']]
```

```

In [32]: fig,axs=plt.subplots(nrows=11,figsize=(12,50))
forecast_base_bill_ele=pd.DataFrame({'Retention':train[train['churn']==0][
    'forecast_base_bill_ele'],
    'Churn':train[train['churn']==1]['forecast_base_bi
ll_ele']})
forecast_base_bill_ele[['Retention','Churn']].plot(kind='hist',bins=50,a
x=axes[0],stacked=True);
axes[0].set_xlabel('forecast_base_bill_ele')
axes[0].ticklabel_format(style='plain',axis='x')
forecast_base_bill_year=pd.DataFrame({'Retention':train[train['churn']==
    0]['forecast_base_bill_year'],
    'Churn':train[train['churn']==1]['forecast_base_bi
ll_year']})
forecast_base_bill_year[['Retention','Churn']].plot(kind='hist',bins=50,
ax=axes[1],stacked=True);
axes[1].set_xlabel('forecast_base_bill_year')
axes[1].ticklabel_format(style='plain',axis='x')
forecast_bill_12m=pd.DataFrame({'Retention':train[train['churn']==0]['fo
recast_bill_12m'],
    'Churn':train[train['churn']==1]['forecast_bill_12
m']})
forecast_bill_12m[['Retention','Churn']].plot(kind='hist',bins=50,ax=axes
[2],stacked=True);
axes[2].set_xlabel('forecast_bill_12m')
axes[2].ticklabel_format(style='plain',axis='x')
forecast_cons=pd.DataFrame({'Retention':train[train['churn']==0]['foreca
st_cons'],
    'Churn':train[train['churn']==1]['forecast_cons'
]})
forecast_cons[['Retention','Churn']].plot(kind='hist',bins=50,ax=axes[3],
stacked=True);
axes[3].set_xlabel('forecast_cons')
axes[3].ticklabel_format(style='plain',axis='x')

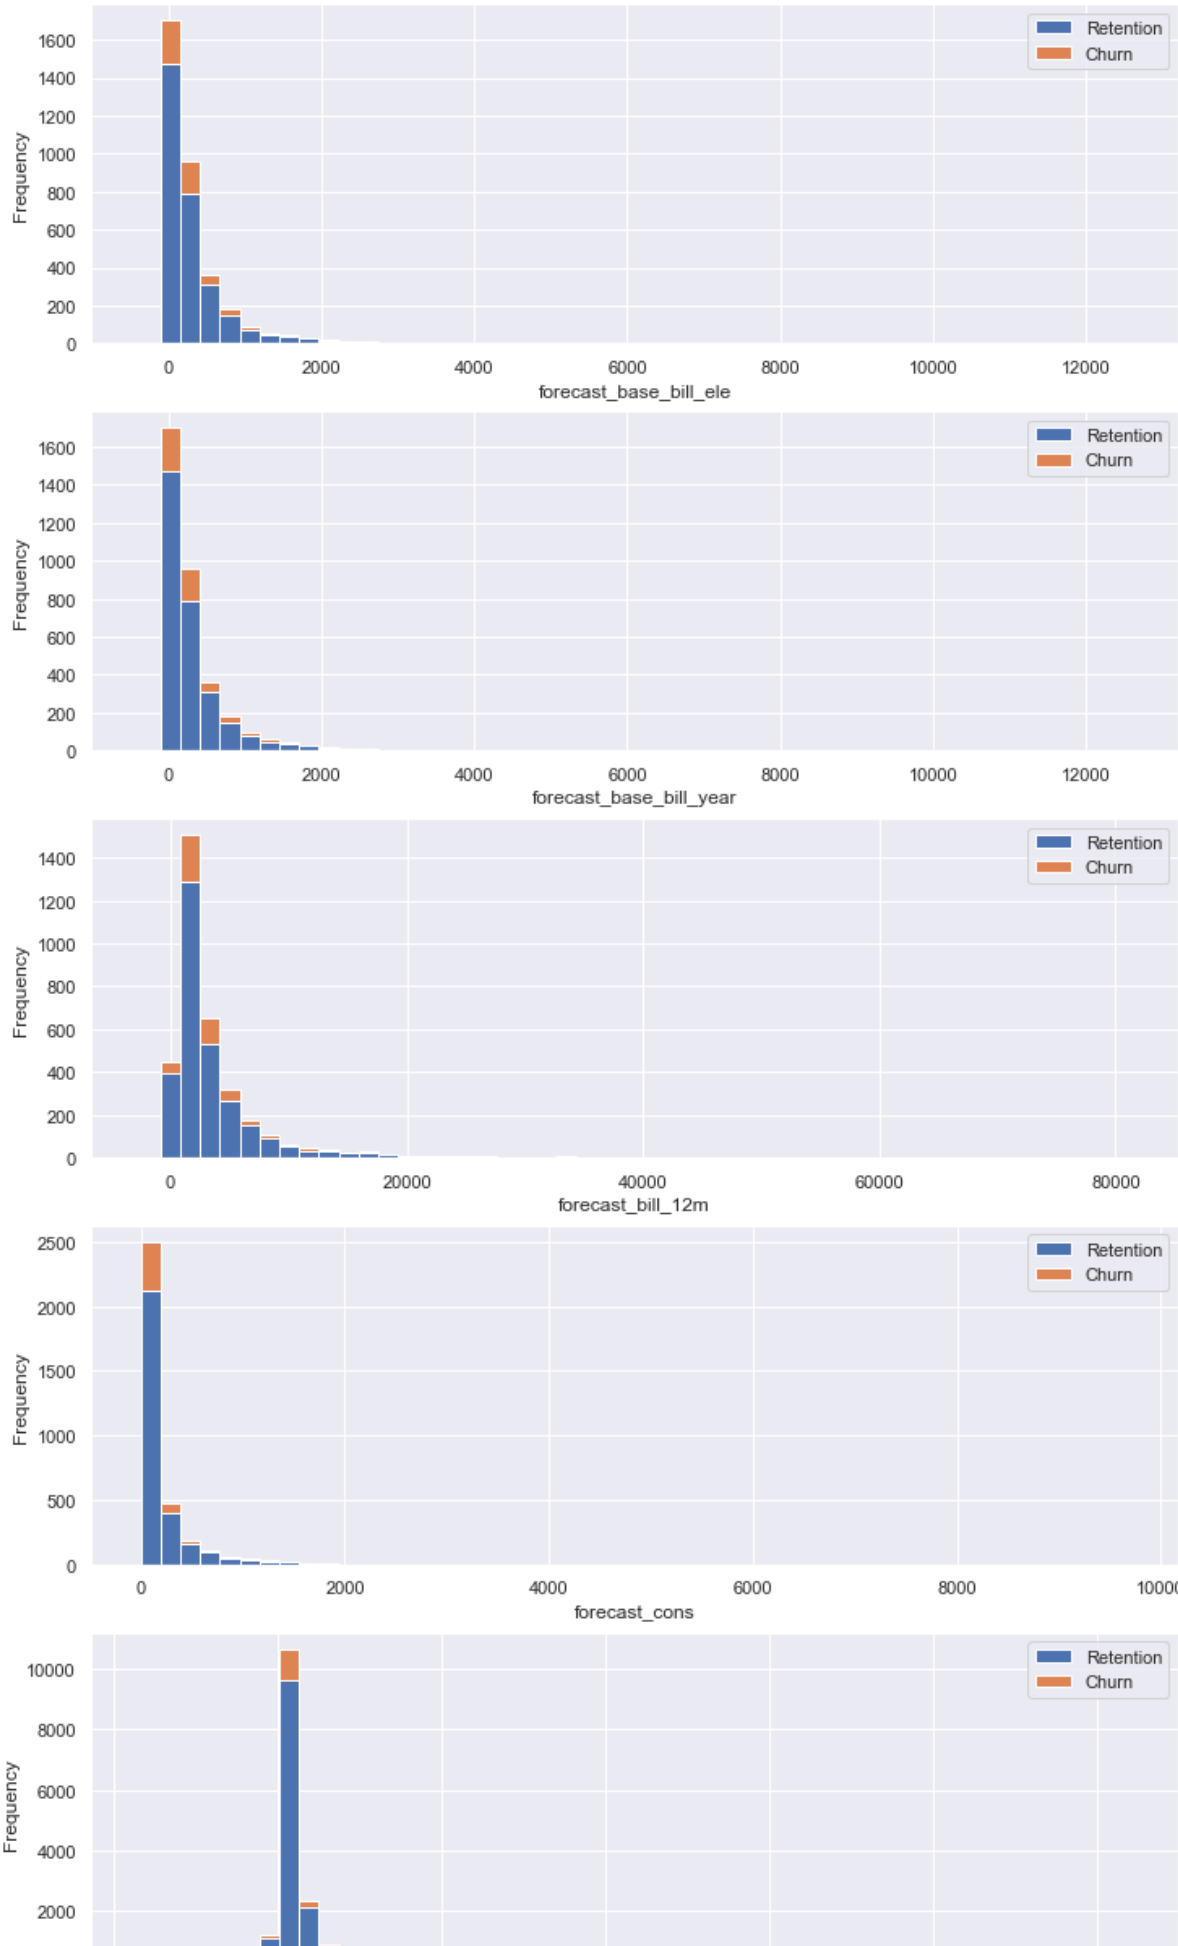
forecast_cons_12m=pd.DataFrame({'Retention':train[train['churn']==0]['fo
recast_cons_12m'],
    'Churn':train[train['churn']==1]['forecast_cons_12
m']})
forecast_cons_12m[['Retention','Churn']].plot(kind='hist',bins=50,ax=axes
[4],stacked=True);
axes[4].set_xlabel('forecast_cons_12m')
axes[4].ticklabel_format(style='plain',axis='x')
forecast_cons_year=pd.DataFrame({'Retention':train[train['churn']==0]['f
orecast_cons_year'],
    'Churn':train[train['churn']==1]['forecast_cons_ye
ar']})
forecast_cons_year[['Retention','Churn']].plot(kind='hist',bins=50,ax=ax
s[5],stacked=True);
axes[5].set_xlabel('forecast_cons_year')
axes[5].ticklabel_format(style='plain',axis='x')
forecast_discount_energy=pd.DataFrame({'Retention':train[train['churn']=
    =0]['forecast_discount_energy'],
    'Churn':train[train['churn']==1]['forecast_discoun
t_energy']})
forecast_discount_energy[['Retention','Churn']].plot(kind='hist',bins=50
,ax=axes[6],stacked=True);

```

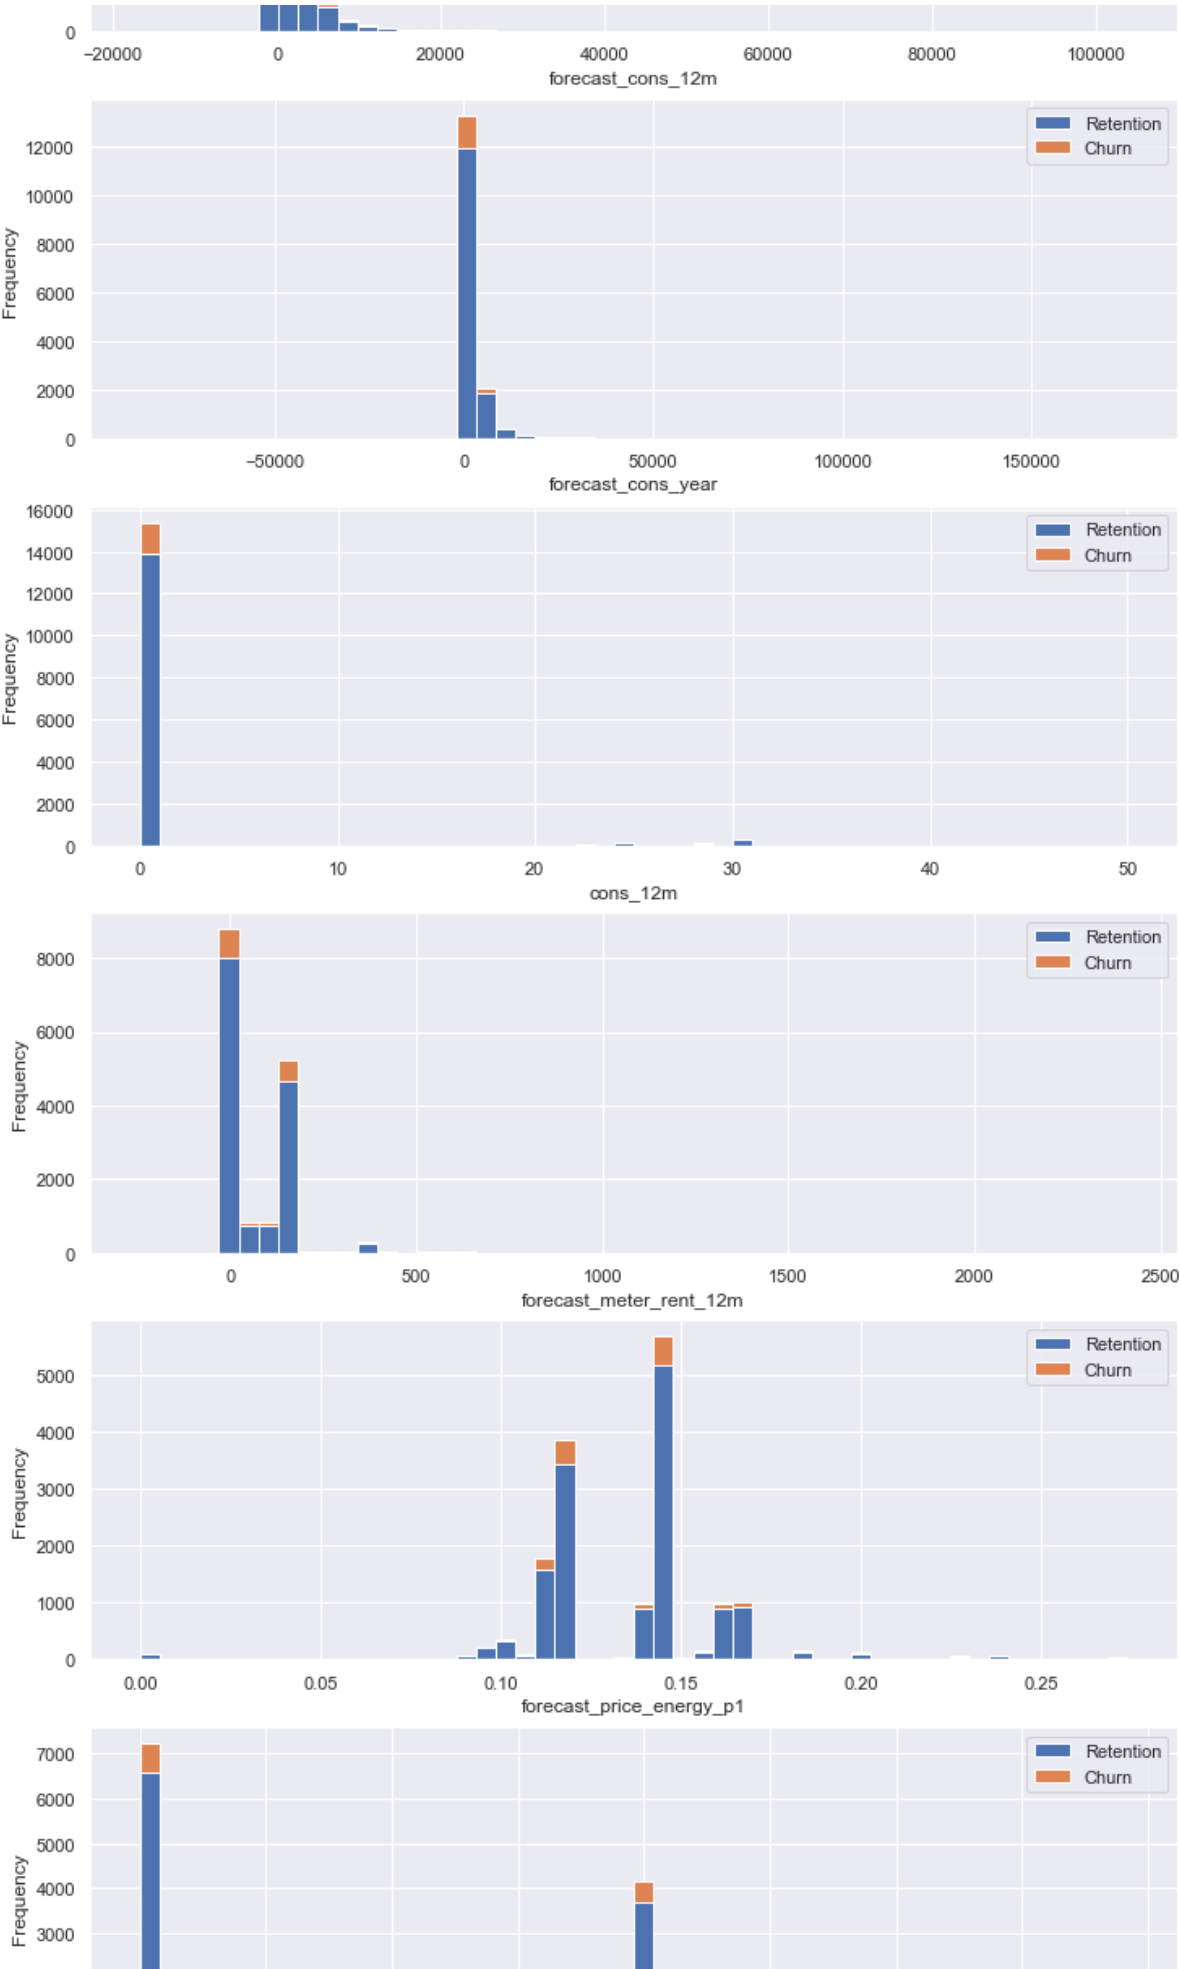
```


axs[6].set_xlabel('cons_12m')
axs[6].ticklabel_format(style='plain',axis='x')
forecast_meter_rent_12m=pd.DataFrame({'Retention':train[train['churn']==
0]['forecast_meter_rent_12m'],
                                     'Churn':train[train['churn']==1]['forecast_meter_r
ent_12m']})
forecast_meter_rent_12m[['Retention','Churn']].plot(kind='hist',bins=50,
ax=axs[7],stacked=True);
axs[7].set_xlabel('forecast_meter_rent_12m')
axs[7].ticklabel_format(style='plain',axis='x')
forecast_price_energy_p1=pd.DataFrame({'Retention':train[train['churn']=
=0]['forecast_price_energy_p1'],
                                     'Churn':train[train['churn']==1]['forecast_price_e
nergy_p1']})
forecast_price_energy_p1[['Retention','Churn']].plot(kind='hist',bins=50
,ax=axs[8],stacked=True);
axs[8].set_xlabel('forecast_price_energy_p1')
axs[8].ticklabel_format(style='plain',axis='x')
forecast_price_energy_p2=pd.DataFrame({'Retention':train[train['churn']=
=0]['forecast_price_energy_p2'],
                                     'Churn':train[train['churn']==1]['forecast_price_e
nergy_p2']})
forecast_price_energy_p2[['Retention','Churn']].plot(kind='hist',bins=50
,ax=axs[9],stacked=True);
axs[9].set_xlabel('forecast_price_energy_p2')
axs[9].ticklabel_format(style='plain',axis='x')
forecast_price_pow_p1=pd.DataFrame({'Retention':train[train['churn']==0]
['forecast_price_pow_p1'],
                                     'Churn':train[train['churn']==1]['forecast_price_p
ow_p1']})
forecast_price_pow_p1[['Retention','Churn']].plot(kind='hist',bins=50,ax
=axs[10],stacked=True);
axs[10].set_xlabel('forecast_price_pow_p1')
axs[10].ticklabel_format(style='plain',axis='x')

```





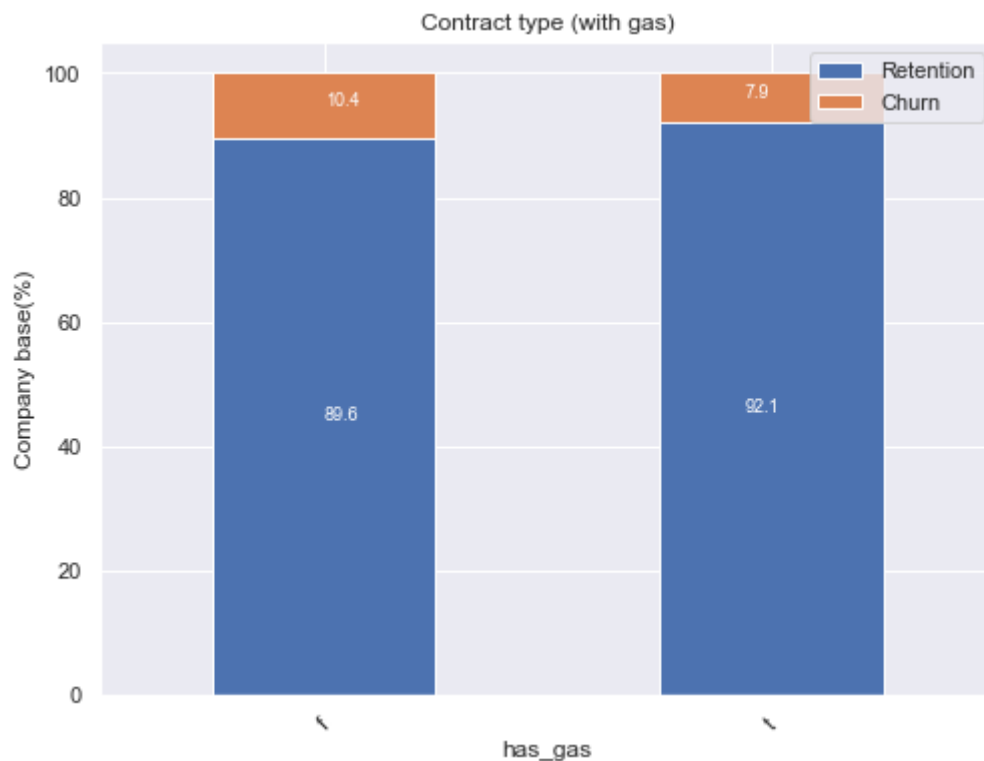




It's similarly to the consumption plots, that lots of variables are highly skewed to the right.

```
In [33]: #Now is for the contract type(electricity,gas)
contract_type=train[['id','has_gas','churn']]
contract=contract_type.groupby([contract_type['churn'],
                                contract_type['has_gas']])['id'].count().
unstack(level=0)
```

```
In [34]: contract_percentage=(contract.div(contract.sum(axis=1),axis=0)*100).sort_
_values(by=[1],ascending=False)
ax=contract_percentage.plot(kind='bar',stacked=True,figsize=(8,6),rot=45
)
for p in ax.patches:
    value=str(round(p.get_height(),1))
    if value=='0':
        continue
    ax.annotate(value,((p.get_x()+p.get_width()/2)*0.94,p.get_y()+p.get_
height()/2*0.99),
                color='white',size=(9))
plt.title('Contract type (with gas)')
plt.legend(['Retention','Churn'],loc="upper right")
plt.ylabel("Company base(%)");
```



```
In [35]: #Now is about Margins
margin=train[['id','margin_gross_pow_ele','margin_net_pow_ele','net_marg
in']]
```

```
In [36]: fig,axs=plt.subplots(nrows=3,figsize=(10,15))
sns.boxplot(margin['margin_gross_pow_ele'],ax=axs[0])
sns.boxplot(margin['margin_net_pow_ele'],ax=axs[1])
sns.boxplot(margin['net_margin'],ax=axs[2])
for ax in axs:
    ax.ticklabel_format(style='plain',axis='x')
plt.show()
```

```
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```



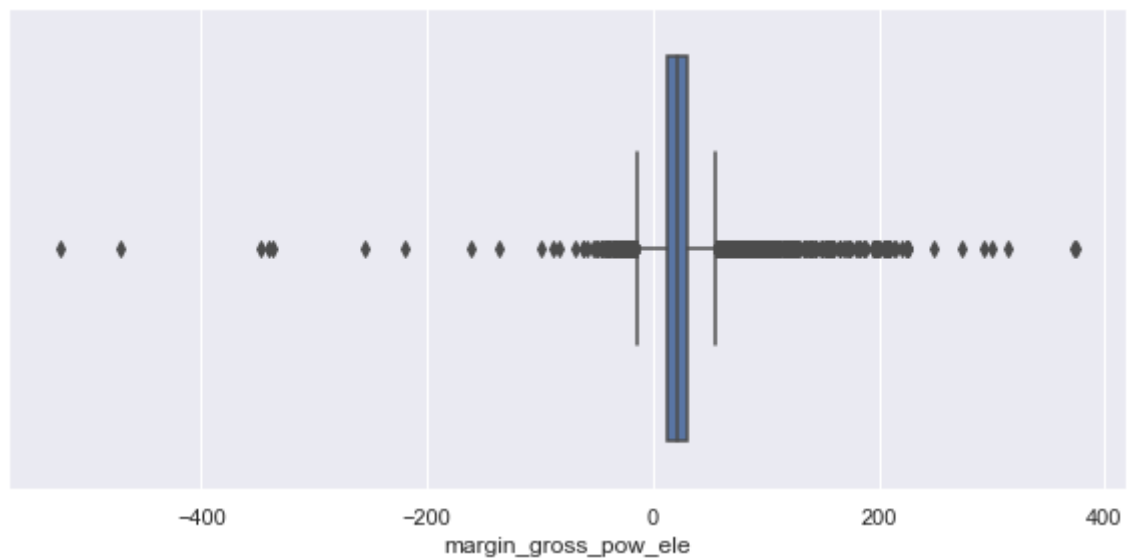
```
warnings.warn(
```

```
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
warnings.warn(
```

```
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```


```
warnings.warn(
```



```
In [37]: #Next is about the Subscribed power  
power=train[['id','pow_max','churn']].fillna(0)
```


```
In [38]: figure=plt.figure()
pow_max=pd.DataFrame({'Retention':power[power['churn']==0]['pow_max'],
                      'Churn':power[power['churn']==1]['pow_max']})
pow_max[['Retention','Churn']].plot(kind='hist',bins=50,stacked=True);
plt.xlabel('pow_max')
plt.ticklabel_format(style='plain',axis='x');
```

<Figure size 432x288 with 0 Axes>



```
In [39]: #Last id for others variables
others=train[['id','nb_prod_act','num_years_antig','origin_up','churn']]
products=others.groupby([others['nb_prod_act'],others['churn']])['id'].count().unstack(level=1)
products_percentage=(products.div(products.sum(axis=1),axis=0)*100).sort_values(by=[1],ascending=False)
```


```
In [40]: ax=products_percentage.plot(kind='bar',stacked=True,figsize=(15,10),rot=
0)
for p in ax.patches:
    value=str(round(p.get_height(),1))
    if value=='0':
        continue
    ax.annotate(value,((p.get_x()+p.get_width()/2)*0.99,p.get_y()+p.get_
height()/2*0.9),
                color='white',size=(9))
plt.title('Number of products')
plt.legend(['Retention','Churn'],loc="upper right")
plt.ylabel("Company base(%)");
```



```
In [41]: years_antig=others.groupby([others['num_years_antig'],others['churn']])[
'id'].count().unstack(level=1)
years_antig_percentage=(years_antig.div(years_antig.sum(axis=1),axis=0)*
100)
```



```
In [42]: ax=years_antig_percentage.plot(kind='bar',stacked=True,figsize=(15,10),rot=0)
for p in ax.patches:
    value=str(round(p.get_height(),1))
    if value=='0':
        continue
    ax.annotate(value,((p.get_x()+p.get_width()/2)*0.99,p.get_y()+p.get_height()/2*0.9),
                color='white',size=(9))
plt.title('Number years')
plt.legend(['Retention','Churn'],loc="upper right")
plt.ylabel("Company base(%)");
```




```
In [43]: origin=others.groupby([others['origin_up'],others['churn']])[ 'id' ].count()
          .unstack(level=1)
origin_percentage=(origin.div(origin.sum(axis=1),axis=0)*100).sort_values(
by=[1],ascending=False)
```

```


In [44]: ax=origin_percentage.plot(kind='bar',stacked=True,figsize=(20,12),rot=0)
for p in ax.patches:
    value=str(round(p.get_height(),1))
    if value=='0':
        continue
    ax.annotate(value,((p.get_x()+p.get_width()/2)*0.99,p.get_y()+p.get_
height()/2*0.9),
                color='white',size=(9))
plt.title('Origin contract/offer')
plt.legend(['Retention','Churn'],loc="upper right")
plt.ylabel("Company base(%)");

```



## Data Cleaning

```
In [45]: #plot the missing data
plt.figure(figsize=(15,12))
(train.isnull().sum()/len(train.index)*100).plot(kind='bar')
plt.xlabel('Variables')
plt.ylabel('Missing values(%)')
plt.show()
```



From the figure above, we can remove the variables that more than 60% values missing

```
In [46]: train.drop(columns=['campaign_disc_ele', 'date_first_activ', 'forecast_base_bill_ele', 'forecast_base_bill_year', 'forecast_bill_12m', 'forecast_cons', 'activity_new'], inplace=True)
```

```
In [47]: #Check The removed dataframe
pd.DataFrame({'Dataframe columns':train.columns})
```

Out[47]:

Dataframe columns	
0	id
1	channel_sales
2	cons_12m
3	cons_gas_12m
4	cons_last_month
5	date_activ
6	date_end
7	date_modif_prod
8	date_renewal
9	forecast_cons_12m
10	forecast_cons_year
11	forecast_discount_energy
12	forecast_meter_rent_12m
13	forecast_price_energy_p1
14	forecast_price_energy_p2
15	forecast_price_pow_p1
16	has_gas
17	imp_cons
18	margin_gross_pow_ele
19	margin_net_pow_ele
20	nb_prod_act
21	net_margin
22	num_years_antig
23	origin_up
24	pow_max
25	churn

```
In [48]: #Check the duplicates
train[train.duplicated()]
```

Out[48]:


id	channel_sales	cons_12m	cons_gas_12m	cons_last_month	date_activ	date_end	date_modi
----	---------------	----------	--------------	-----------------	------------	----------	-----------

0 rows × 26 columns

There seems no duplicated data of the train dataframe

```
In [49]: #Check the history missing data
missing_data_percentage=history_data.isnull().sum()/len(history_data.index)*100
```

```
In [50]: plt.figure(figsize=(15,12))
missing_data_percentage.plot(kind='bar')
plt.xlabel('Variables')
plt.ylabel('Missing values(%)')
plt.show()
```



There is not much data missing, we will substitute them with the median in the next step

## Formating data

```
In [51]: #fill the missing date with the median date which use the value_counts()
train.loc[train['date_modif_prod'].isnull(), 'date_modif_prod'] = train['date_modif_prod'].value_counts().index[0]
train.loc[train['date_end'].isnull(), 'date_end'] = train['date_end'].value_counts().index[0]
train.loc[train['date_renewal'].isnull(), 'date_renewal'] = train['date_renewal'].value_counts().index[0]
```

```
In [52]: #fill the price data with median
history_data.loc[history_data['price_p1_var'].isnull(), 'price_p1_var'] = history_data['price_p1_var'].median()
history_data.loc[history_data['price_p2_var'].isnull(), 'price_p2_var'] = history_data['price_p2_var'].median()
history_data.loc[history_data['price_p3_var'].isnull(), 'price_p3_var'] = history_data['price_p3_var'].median()
history_data.loc[history_data['price_p1_fix'].isnull(), 'price_p1_fix'] = history_data['price_p1_fix'].median()
history_data.loc[history_data['price_p2_fix'].isnull(), 'price_p2_fix'] = history_data['price_p2_fix'].median()
history_data.loc[history_data['price_p3_fix'].isnull(), 'price_p3_fix'] = history_data['price_p3_fix'].median()
```

```
In [53]: #fill the negative data of history with median
history_data.loc[history_data['price_p1_fix'] < 0, 'price_p1_fix'] = history_data['price_p1_fix'].median()
history_data.loc[history_data['price_p2_fix'] < 0, 'price_p2_fix'] = history_data['price_p2_fix'].median()
history_data.loc[history_data['price_p3_fix'] < 0, 'price_p3_fix'] = history_data['price_p3_fix'].median()
```

```
In [54]: #Transform date columns to datetime type
train['date_activ'] = pd.to_datetime(train['date_activ'], format='%Y-%m-%d')
train['date_end'] = pd.to_datetime(train['date_end'], format='%Y-%m-%d')
train['date_modif_prod'] = pd.to_datetime(train['date_modif_prod'], format='%Y-%m-%d')
train['date_renewal'] = pd.to_datetime(train['date_renewal'], format='%Y-%m-%d')
history_data['price_date'] = pd.to_datetime(history_data['price_date'], format='%Y-%m-%d')
```

## Saving data

```
In [55]: #Make directly processed data if it does not exist
train.to_csv('train_clean.csv', index = False)
history_data.to_csv('history_clean.csv', index = False)
```

In [ ]:

In [ ]:

In [ ]: